

CS 415 Parallel Computing

Mandelbrot Sets

March 2, 2012

Erin Keith

Introduction:

This project focused on computing Mandelbrot sets and outputting an image which illustrates values contained in that set. A Mandelbrot set is a set of complex numbers, such that

$$M = \left\{ c \in \mathbb{C} \mid \lim_{n \rightarrow \infty} Z_n \neq \infty \right\}$$

where:

$$\begin{aligned} Z_0 &= c \\ Z_{n+1} &= Z_n^2 + c \end{aligned}$$

Since the complex set can be represented in a two dimensional Cartesian coordinate system, and we can limit the number of iterations to 256, the form in Figure 1 will be displayed.




Figure 1: Mandelbrot Fractal

Procedure:

Sequential Program

To be able to determine speedup time, it was imperative to first create the algorithm, as provided in the book, sequentially. A struct, containing real and imaginary variables, was helpful for handling complex numbers. One complex struct handled the Cartesian coordinates as a representation of the set of complex numbers between $-2-2i$, and $2+2i$, scaled by an image factor. Another complex struct started at zero and throughout the iterations the following formulas were computed:

$$z_{\text{real}} = z_{\text{real}}^2 - z_{\text{imag}}^2 + c_{\text{real}}$$

$$z_{\text{imag}} = 2z_{\text{real}}z_{\text{imag}} + c_{\text{imag}}$$

as long as the magnitude of z was less than 2 and the iterations were less than 256. The number of iterations was then stored as the pixel value for the image at that coordinate.

To determine the amount of time it took to calculate images of different sizes, each image size runtime was averaged over 10 executions on the grid, and the runtime values were gathered in Table 1

Table 1: Sequential Time over Pixels

| Pixels Dimensions | Time (sec) |
|-------------------|------------|
| 500 | 0.126 |
| 1000 | 0.253 |
| 1500 | 0.485 |
| 2000 | 0.798 |
| 2500 | 1.226 |
| 3000 | 1.793 |
| 3500 | 2.394 |
| 4000 | 3.052 |
| 4500 | 3.882 |
| 5000 | 4.817 |
| 5500 | 5.801 |
| 6000 | 6.828 |
| 6500 | 8.127 |
| 7000 | 10.15 |
| 7500 | 12.126 |
| 8000 | 16.479 |
| 8500 | 18.824 |
| 9000 | 22.705 |
| 9500 | 26.058 |
| 10000 | 35.256 |

The data gathered in Table 1 is represented by Figure 2, which clearly shows a curve resembling an exponential growth trend.

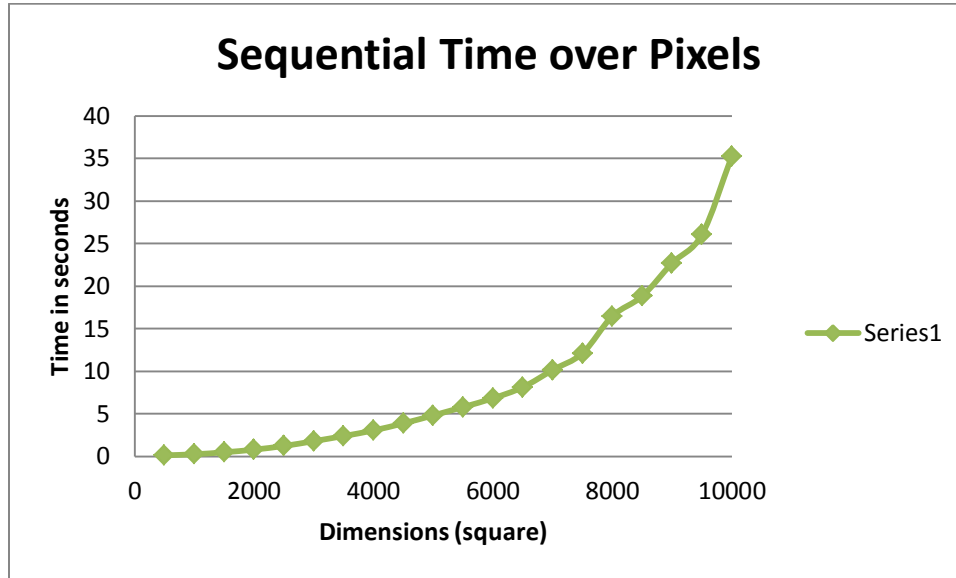


Figure 2: Sequential Runtime for Different Image Sizes

Part II: Static Parallel

The next part of the assignment was to design a parallel algorithm, assigning the slave processes a fixed range of rows for which to compute the values. Since the dimensions of my image were declared as global constants, I determined that it would be redundant to send the slaves their ranges. I made the slaves “intelligent” and had them compute their range to save communication time. The master then determined the appropriate range based on the ID of the process whose message it received. Runtimes for each image dimension was run on a range of processors, from 2 to 20 in two processor increments; the data is displayed in Table 2.

Table 2: Static Parallel Runtimes over Pixels and Processors

| Pixels | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 500 | 0.073 | 0.064 | 0.049 | 0.041 | 0.038 | 0.036 | 0.036 | 0.036 | 0.033 | 0.029 |
| 1000 | 0.202 | 0.167 | 0.121 | 0.125 | 0.106 | 0.085 | 0.074 | 0.067 | 0.100 | 0.191 |
| 1500 | 0.577 | 0.501 | 0.367 | 0.280 | 0.230 | 0.189 | 0.256 | 0.304 | 0.271 | 0.241 |
| 2000 | 1.027 | 0.890 | 0.652 | 0.500 | 0.408 | 0.337 | 0.459 | 0.538 | 0.523 | 0.408 |
| 2500 | 1.610 | 1.399 | 1.015 | 0.779 | 0.634 | 0.525 | 0.455 | 0.329 | 0.281 | 0.284 |
| 3000 | 2.323 | 2.005 | 1.458 | 1.119 | 0.913 | 0.751 | 0.648 | 0.470 | 0.616 | 0.826 |
| 3500 | 3.159 | 2.5 | 1.985 | 1.518 | 1.241 | 1.025 | 1.296 | 1.265 | 1.294 | 1.221 |
| 4000 | 4.110 | 3.0 | 2.5 | 1.987 | 1.629 | 1.335 | 1.530 | 1.772 | 1.746 | 1.442 |
| 4500 | 5.205 | 4.0 | 3.0 | 2.512 | 2.040 | 1.692 | 1.564 | 1.615 | 1.515 | 1.294 |
| 5000 | 6.975 | 5.342 | 3.800 | 3.0 | 2.886 | 2.284 | 1.545 | 1.282 | 1.172 | 1.434 |
| 5500 | 6.091 | 5.041 | 3.366 | 3.0 | 2.5 | 2.051 | 2.121 | 1.5 | 1.4 | 2.0 |

| | | | | | | | | | | |
|------|--------|--------|-------|-------|-------|-----|-------|-------|-------|-------|
| 6000 | 7.289 | 6.000 | 4.346 | 3.365 | 2.725 | 2.5 | 1.936 | 1.761 | 1.878 | 2.228 |
| 6500 | 8.539 | 7.043 | 5.097 | 3.952 | 3.202 | 3.0 | 2.5 | 2.110 | 2.198 | 2.458 |
| 7000 | 9.905 | 8.296 | 5.904 | 4.574 | 3.706 | 3.5 | 3.0 | 2.5 | 2.439 | 2.649 |
| 7500 | 12.098 | 10.617 | 7.707 | 7.0 | 6.0 | 5.0 | 4.0 | 3.0 | 3.0 | 3.0 |

Since this was the second time running the data, with a much larger data set, I found an error in my program, which creates a segmentation fault for predictable conditions. Unfortunately, I did not have time to address this problem, so only a range of 500 to 7500 square pixel images was used. The missing data was filled in with round numbers (indicated by two decimal places) to fill in gaps in the graph. Figure 3 displays the graphical representation of the runtime data from Table 2.




Figure 3: Static Parallel Runtimes

Part III: Dynamic Parallel

Finally, I attacked the problem of dynamically assigning rows from a workpool to processors as they became available. Instead of using the book's suggestion to send the array of pixel values and the row number, I used a lookup table to keep track of the rows being processed based on processor ID, which helped cut down on communication time.

The runtimes created by the dynamic parallel program are contained in Table 3.

Table 3: Dynamic Parallel Runtimes over Pixels and Processors

| Pixels | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 500 | 0.067 | 0.046 | 0.038 | 0.040 | 0.031 | 0.028 | 0.028 | 0.027 | 0.104 | 0.103 |
| 1000 | 0.253 | 0.123 | 0.099 | 0.084 | 0.074 | 0.069 | 0.064 | 0.062 | 0.139 | 0.137 |

| | | | | | | | | | | |
|-------|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1500 | 0.570 | 0.259 | 0.178 | 0.151 | 0.147 | 0.138 | 0.125 | 0.129 | 0.176 | 0.194 |
| 2000 | 1.015 | 0.453 | 0.359 | 0.238 | 0.207 | 0.182 | 0.166 | 0.188 | 0.264 | 0.241 |
| 2500 | 2.372 | 1.289 | 0.839 | 0.567 | 0.415 | 0.339 | 0.345 | 0.324 | 0.380 | 0.406 |
| 3000 | 3.171 | 1.280 | 0.683 | 0.544 | 0.471 | 0.478 | 0.438 | 0.434 | 0.500 | 0.487 |
| 3500 | 3.103 | 1.263 | 0.835 | 0.406 | 0.369 | 0.562 | 0.544 | 0.526 | 0.540 | 0.580 |
| 4000 | 4.091 | 2.029 | 1.498 | 1.226 | 1.056 | 0.958 | 0.666 | 0.651 | 0.715 | 0.697 |
| 4500 | 5.121 | 2.059 | 1.362 | 1.155 | 0.971 | 0.941 | 0.888 | 0.839 | 0.834 | 0.797 |
| 5000 | 4.940 | 2.735 | 1.929 | 1.391 | 1.221 | 1.099 | 0.987 | 0.979 | 1.000 | 0.959 |
| 5500 | 5.894 | 2.721 | 1.848 | 2.282 | 1.704 | 1.891 | 1.521 | 1.433 | 1.518 | 1.625 |
| 6000 | 8.111 | 3.041 | 2.318 | 1.834 | 1.691 | 1.585 | 1.419 | 1.235 | 1.208 | 1.128 |
| 6500 | 10.622 | 4.199 | 2.955 | 2.259 | 1.874 | 1.617 | 1.497 | 1.391 | 1.845 | 1.892 |
| 7000 | 12.398 | 4.746 | 3.290 | 2.268 | 2.498 | 2.242 | 2.064 | 2.007 | 1.520 | 1.295 |
| 7500 | 14.245 | 5.620 | 3.714 | 3.170 | 2.675 | 2.064 | 2.815 | 2.565 | 2.563 | 2.628 |
| 8000 | 16.210 | 6.336 | 4.180 | 3.951 | 3.669 | 3.280 | 3.004 | 3.049 | 3.022 | 2.979 |
| 8500 | 16.248 | 6.931 | 4.783 | 4.018 | 3.681 | 3.138 | 3.048 | 2.826 | 2.841 | 2.928 |
| 9000 | 24.080 | 8.723 | 5.989 | 5.160 | 4.492 | 3.949 | 3.820 | 3.594 | 3.412 | 3.574 |
| 9500 | 22.891 | 8.911 | 5.910 | 5.266 | 4.669 | 4.702 | 4.373 | 4.442 | 4.479 | 4.362 |
| 10000 | 25.159 | 8.910 | 5.888 | 4.989 | 4.449 | 4.112 | 3.998 | 3.806 | 3.928 | 3.888 |

Figure 4 illustrates the runtime of the dynamic parallel program over the same ranges as the static parallel program.




Figure 4: Dynamic Parallel Runtimes

Analysis:

The curve generated by the sequential program clearly demonstrates some sort of threshold around the 8000 x 8000 pixel images. Otherwise the curve shows the expected linear increase.

The statically parallel implementation of the program shows a shallower, more linear curve than the dynamic, as I'm dividing the work by processor, directly. It is important to note that the dynamic program's overall runtimes were less than the static program, except for the runtimes for two processors, which were actually faster in the static program, since there was significantly less communication.

The overall runtimes of the dynamically implemented parallel program show a more sharply decreasing trend for increasing processors than the static program. Using two processors for the parallel programs simulates the sequential environment. There is a steep drop from two to four processors, as this is actually the first instance of multiple processors being used. If you remove the two processor data, a trough is evident, on the graph. If I had had enough time, I would have analyzed these minimums to try to determine some sort of relationship, which would lead to an optimal percentage of work to dole out to each processor.

There is also significant speedup of parallel over sequential algorithms, as given by the following equation:

$$\text{Speedup} = \frac{t_{\text{sequential}}}{t_{\text{dynamic parallel}}}$$

The predictably increasing trend, with a maximum speedup of roughly 9, is contained in Table 4.

Table 4: Speedup of Dynamic Parallel Runtimes over Sequential

| Pixels | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 500 | 1.867 | 2.739 | 3.290 | 3.179 | 4.090 | 4.436 | 4.572 | 4.716 | 1.209 | 1.223 |
| 1000 | 0.998 | 2.065 | 2.546 | 3.028 | 3.403 | 3.641 | 3.934 | 4.100 | 1.818 | 1.848 |
| 1500 | 0.851 | 1.876 | 2.723 | 3.215 | 3.298 | 3.511 | 3.888 | 3.753 | 2.755 | 2.496 |
| 2000 | 0.786 | 1.763 | 2.226 | 3.351 | 3.855 | 4.391 | 4.811 | 4.254 | 3.023 | 3.306 |
| 2500 | 0.517 | 0.951 | 1.462 | 2.164 | 2.958 | 3.619 | 3.553 | 3.778 | 3.230 | 3.016 |
| 3000 | 0.566 | 1.401 | 2.626 | 3.295 | 3.805 | 3.755 | 4.092 | 4.135 | 3.585 | 3.679 |
| 3500 | 0.772 | 1.895 | 2.868 | 5.890 | 6.492 | 4.259 | 4.398 | 4.548 | 4.437 | 4.129 |
| 4000 | 0.746 | 1.504 | 2.037 | 2.489 | 2.891 | 3.187 | 4.582 | 4.692 | 4.269 | 4.377 |
| 4500 | 0.758 | 1.885 | 2.851 | 3.361 | 4.000 | 4.124 | 4.371 | 4.629 | 4.656 | 4.870 |
| 5000 | 0.975 | 1.761 | 2.497 | 3.464 | 3.945 | 4.384 | 4.883 | 4.918 | 4.816 | 5.025 |
| 5500 | 0.984 | 2.132 | 3.138 | 2.542 | 3.404 | 3.068 | 3.815 | 4.048 | 3.822 | 3.570 |
| 6000 | 0.842 | 2.245 | 2.945 | 3.723 | 4.038 | 4.309 | 4.812 | 5.527 | 5.651 | 6.055 |
| 6500 | 0.765 | 1.935 | 2.750 | 3.597 | 4.338 | 5.026 | 5.430 | 5.843 | 4.404 | 4.295 |
| 7000 | 0.819 | 2.139 | 3.085 | 4.475 | 4.064 | 4.526 | 4.918 | 5.058 | 6.677 | 7.839 |
| 7500 | 0.851 | 2.158 | 3.265 | 3.825 | 4.534 | 5.875 | 4.308 | 4.727 | 4.730 | 4.614 |
| 8000 | 1.017 | 2.601 | 3.943 | 4.171 | 4.492 | 5.024 | 5.486 | 5.405 | 5.452 | 5.532 |
| 8500 | 1.159 | 2.716 | 3.935 | 4.685 | 5.114 | 5.998 | 6.177 | 6.662 | 6.627 | 6.430 |

| | | | | | | | | | | |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 9000 | 0.943 | 2.603 | 3.791 | 4.400 | 5.055 | 5.750 | 5.943 | 6.317 | 6.654 | 6.354 |
| 9500 | 1.138 | 2.924 | 4.409 | 4.949 | 5.581 | 5.542 | 5.959 | 5.867 | 5.818 | 5.974 |
| 10000 | 1.401 | 3.957 | 5.988 | 7.067 | 7.924 | 8.574 | 8.819 | 9.262 | 8.976 | 9.068 |

Speedup for each image size and processor is illustrated by Figure 5. Then general trend is increased speedup as image size and the number of processors increase, which is to be expected. The surprise is that the processor number with greatest speedup is 16, for which I cannot think of an explanation (except that maybe the grid was running extra fast).




Figure 5: Speedup of Dynamic Parallel over Sequential

Finally, and most interestingly, is the efficiency of each processor in speeding up processing of different image sizes, shown in Table 5. Here we have:

$$\text{Efficiency} = \frac{\text{Speedup}}{\text{Number of Processors}}$$

Table 5: Efficiency of Dynamic Parallel Runtimes

| Pixels | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 500 | 0.934 | 0.685 | 0.548 | 0.397 | 0.409 | 0.370 | 0.327 | 0.295 | 0.067 | 0.061 |
| 1000 | 0.499 | 0.516 | 0.424 | 0.379 | 0.340 | 0.303 | 0.281 | 0.256 | 0.101 | 0.092 |
| 1500 | 0.426 | 0.469 | 0.454 | 0.402 | 0.330 | 0.293 | 0.278 | 0.235 | 0.153 | 0.125 |
| 2000 | 0.393 | 0.441 | 0.371 | 0.419 | 0.386 | 0.366 | 0.344 | 0.266 | 0.168 | 0.165 |
| 2500 | 0.258 | 0.238 | 0.244 | 0.270 | 0.296 | 0.302 | 0.254 | 0.236 | 0.179 | 0.151 |
| 3000 | 0.283 | 0.350 | 0.438 | 0.412 | 0.380 | 0.313 | 0.292 | 0.258 | 0.199 | 0.184 |
| 3500 | 0.386 | 0.474 | 0.478 | 0.736 | 0.649 | 0.355 | 0.314 | 0.284 | 0.247 | 0.206 |
| 4000 | 0.373 | 0.376 | 0.340 | 0.311 | 0.289 | 0.266 | 0.327 | 0.293 | 0.237 | 0.219 |

| | | | | | | | | | | |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 4500 | 0.379 | 0.471 | 0.475 | 0.420 | 0.400 | 0.344 | 0.312 | 0.289 | 0.259 | 0.243 |
| 5000 | 0.488 | 0.440 | 0.416 | 0.433 | 0.395 | 0.365 | 0.349 | 0.307 | 0.268 | 0.251 |
| 5500 | 0.492 | 0.533 | 0.523 | 0.318 | 0.340 | 0.256 | 0.272 | 0.253 | 0.212 | 0.179 |
| 6000 | 0.421 | 0.561 | 0.491 | 0.465 | 0.404 | 0.359 | 0.344 | 0.345 | 0.314 | 0.303 |
| 6500 | 0.383 | 0.484 | 0.458 | 0.450 | 0.434 | 0.419 | 0.388 | 0.365 | 0.245 | 0.215 |
| 7000 | 0.409 | 0.535 | 0.514 | 0.559 | 0.406 | 0.377 | 0.351 | 0.316 | 0.371 | 0.392 |
| 7500 | 0.426 | 0.539 | 0.544 | 0.478 | 0.453 | 0.490 | 0.308 | 0.295 | 0.263 | 0.231 |
| 8000 | 0.508 | 0.650 | 0.657 | 0.521 | 0.449 | 0.419 | 0.392 | 0.338 | 0.303 | 0.277 |
| 8500 | 0.579 | 0.679 | 0.656 | 0.586 | 0.511 | 0.500 | 0.441 | 0.416 | 0.368 | 0.322 |
| 9000 | 0.471 | 0.651 | 0.632 | 0.550 | 0.506 | 0.479 | 0.424 | 0.395 | 0.370 | 0.318 |
| 9500 | 0.569 | 0.731 | 0.735 | 0.619 | 0.558 | 0.462 | 0.426 | 0.367 | 0.323 | 0.299 |
| 10000 | 0.701 | 0.989 | 0.998 | 0.883 | 0.792 | 0.714 | 0.630 | 0.579 | 0.499 | 0.453 |




Figure 6: Efficiency of Dynamic Parallel Process

As the analysis of dynamic runtimes led to insight of how to better assign workload to processors, I also think that these efficiency percentages hold the key. I would have liked to spend time finding the conditions for best and worst efficiencies, and to work out a relationship between them.

Conclusion:

My Mandelbrot programs were successful in demonstrating the speedup of a parallel implementation over a sequential one. My implementations also provide a good example of increased efficiency in dynamically parallel programs over statically parallel programs. Although I am especially pleased that my designs reduced communication time significantly, there are still many aspects of this program to analyze to produce better efficiencies and runtimes.