

CS 415 Parallel Computing

# Matrix Multiplication

April 18, 2012

Erin Keith

## Introduction:

This project focused on computing the product of two matrices. The program was given two matrices to multiply together by summing the product of the first matrix's row elements and the second matrix column elements, to produce each element of the resulting matrix. A matrix with **n** rows by **p** columns multiplied by a matrix with **p** rows by **m** columns results in a matrix of **n** rows by **m** columns:

$$\begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} \times \begin{matrix} 2 & 2 \\ 2 & 2 \\ 2 & 2 \end{matrix} = \begin{matrix} 6 & 6 \\ 6 & 6 \end{matrix}$$

## Procedure:

### *Sequential Program*

The sequential program was very simply comprised of 3 nested for loops used to calculate the sum for each of the resulting elements.

### *Part II: Parallel*


My approach to parallelizing the problem was very similar to my Mandelbrot implementation. After broadcasting one of the matrices, I allowed for a workpool approach to dole out each row of the remaining matrix. The worker process then calculated the sum of the products of elements and returned the resulting array to the master process. The master process was responsible for inserting the row into the appropriate spot in the result matrix.

## Analysis:

Because the sequential program has a time complexity of  $O(n^3)$ , significant speedup was expected. The runtimes for the sequential version of the program are included in Table 1.

**Table 1: Sequential Runtimes for Matrix Multiplication**

Elements	Time
200	0.29
400	1.41
600	4.33
800	10.66
1000	23.85
1200	40.76
1400	67.16
1600	101.57
1800	132.23
2000	183.74




**Figure 1: Runtimes for Sequential Matrix Multiplication Implementation**

Initially, I experienced some superlinear speedup, as did some of my peers. Since my parallel implementation only had two for loops in the worker, I suggest that the parallel implementation far better computation times. As it turns out, I had been dynamically declaring the arrays in my sequential program but not in my parallel, which accounted for the slower sequential runtimes. Table 2 presents runtimes for different number of processors and matrix sizes, which are visualized in Figure 2.

The drawback of my implementation was the amount of communication I included. I did not implement the look up table for keeping track of rows being processed so I had significantly more communication. I ran my parallel program for a range of processors, but had a hard time getting consistent runtimes. I think most of the runtime problems had to do with network congestion and the inconsistency of running in the common queue.

**Table 2: Parallel Runtimes for Matrix Multiplication**

Elements	10	11	12	13
200	5.03	5.89	5.25	3.99
400	5.86	7.42	5.93	5.2
600	4.58	4.94	4.61	5.58
800	5.04	5.82	6.03	7.67
1000	6.09	6.14	7.34	8.01
1200	7.97	7.27	9.62	8.36
1400	9.52	9.54	10.29	8.6
1600	11.61	13.13	11.69	9.21
1800	16.09	14.58	14.6	11.13
2000	19.21	18.64	17.59	14.89



**Figure 2: Runtimes for Parallel Matrix Multiplication Implementation**

I also noticed that there wasn't much difference with increasing the processors. In fact it appeared that runtimes were increasing due to increased processors. This is most likely an effect of too much communication. The speedup and efficiency graphs shown in Figures 2 and 3 are bumpy, demonstrating the inconsistency of the grid.

**Table 3: Speedup for Matrix Multiplication**

Elements	10	11	12	13
<b>200</b>	0.057654	0.239389	0.055238	0.072682
<b>400</b>	0.240614	0.190027	0.237774	0.271154
<b>600</b>	0.945415	2.157895	0.939262	0.775986
<b>800</b>	2.115079	1.831615	1.767828	1.389831
<b>1000</b>	3.916256	6.638436	3.249319	2.977528
<b>1200</b>	5.114178	5.606602	4.237006	4.875598
<b>1400</b>	7.054622	10.64675	6.526725	7.809302
<b>1600</b>	8.748493	7.73572	8.688623	11.02823
<b>1800</b>	8.218148	9.069273	9.056849	11.8805
<b>2000</b>	9.56481	9.857296	10.44571	12.33983

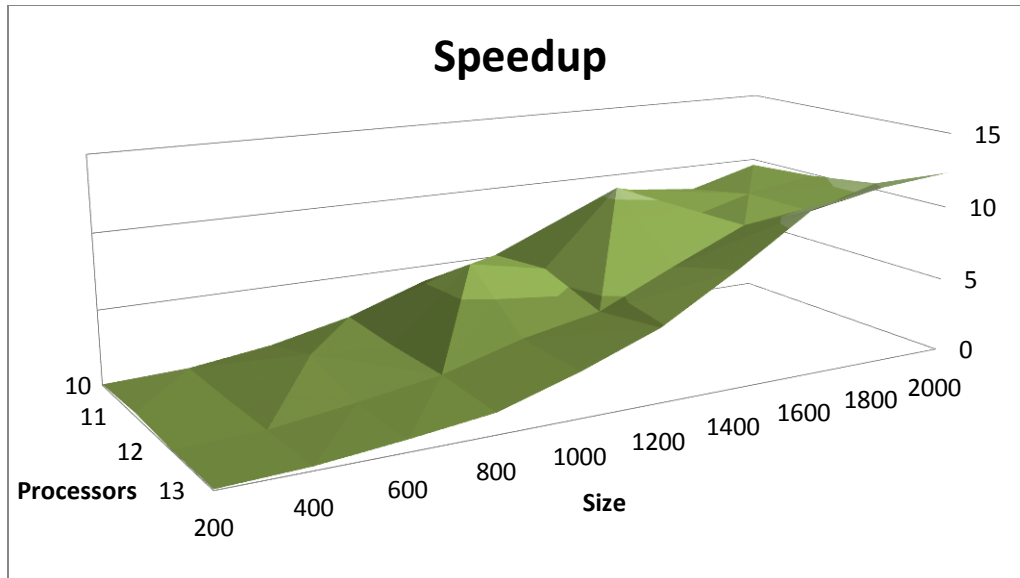


Figure 3: Speedup for Matrix Multiplication Implementation

Table 4: Efficiency for Matrix Multiplication

Elements	10	11	12	13
200	0.576541	2.176262	0.460317	0.55909
400	2.406143	1.727518	1.98145	2.085799
600	9.454148	19.61722	7.827187	5.96912
800	21.15079	16.65105	14.7319	10.691
1000	39.16256	60.34942	27.07766	22.90406
1200	51.14178	50.96911	35.30839	37.5046
1400	70.54622	96.78864	54.38937	60.07156
1600	87.48493	70.32472	72.40519	84.83254
1800	82.18148	82.44794	75.47374	91.38849
2000	95.6481	89.61178	87.04756	94.92173

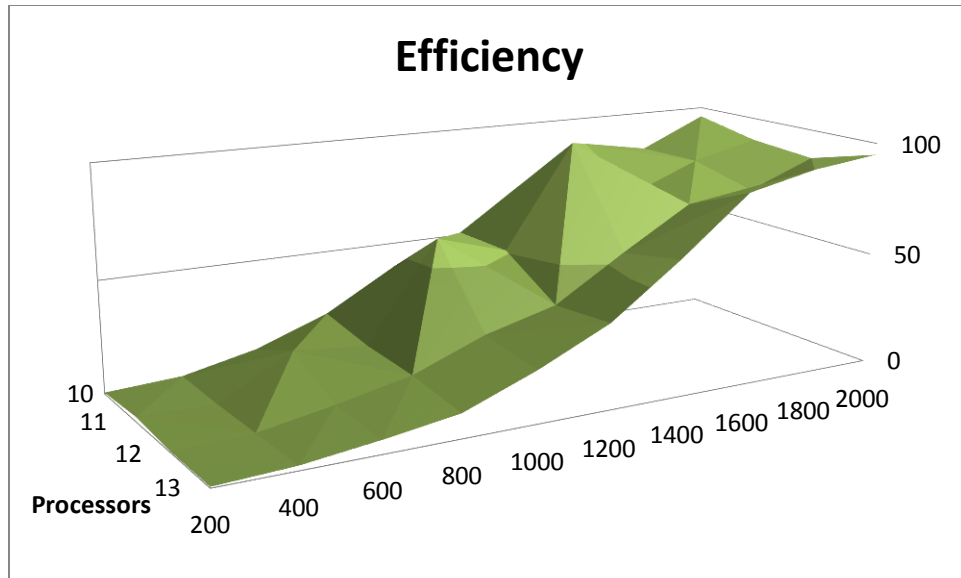


Figure 4: Efficiency for Matrix Multiplication Implementation

**Conclusion:**

This parallel implementation was very similar to that of Mandelbrot, with the addition of utilizing the MPI Broadcast function. I, along with the other students I spoke with, was pretty uncomfortable with efficiency greater than 100 percent, as we should have been. This was a good lesson in taking a closer look at differences between implementations.